

Wine Quality Prediction

July 22, 2023

```
[35]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sb

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from sklearn.svm import SVC
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression

import warnings
warnings.filterwarnings('ignore')
```

```
[34]: df = pd.read_csv('winequalityprediction-red.csv')
print(df.head())
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
[4]: df.describe().T
```

```
[4]:
```

	count	mean	std	min	25%	\
fixed acidity	1599.0	8.319637	1.741096	4.60000	7.1000	
volatile acidity	1599.0	0.527821	0.179060	0.12000	0.3900	
citric acid	1599.0	0.270976	0.194801	0.00000	0.0900	
residual sugar	1599.0	2.538806	1.409928	0.90000	1.9000	
chlorides	1599.0	0.087467	0.047065	0.01200	0.0700	
free sulfur dioxide	1599.0	15.874922	10.460157	1.00000	7.0000	
total sulfur dioxide	1599.0	46.467792	32.895324	6.00000	22.0000	
density	1599.0	0.996747	0.001887	0.99007	0.9956	
pH	1599.0	3.311113	0.154386	2.74000	3.2100	
sulphates	1599.0	0.658149	0.169507	0.33000	0.5500	
alcohol	1599.0	10.422983	1.065668	8.40000	9.5000	
quality	1599.0	5.636023	0.807569	3.00000	5.0000	

	50%	75%	max
fixed acidity	7.90000	9.200000	15.90000
volatile acidity	0.52000	0.640000	1.58000
citric acid	0.26000	0.420000	1.00000
residual sugar	2.20000	2.600000	15.50000
chlorides	0.07900	0.090000	0.61100
free sulfur dioxide	14.00000	21.000000	72.00000
total sulfur dioxide	38.00000	62.000000	289.00000
density	0.99675	0.997835	1.00369
pH	3.31000	3.400000	4.01000

sulphates	0.62000	0.730000	2.00000
alcohol	10.20000	11.100000	14.90000
quality	6.00000	6.000000	8.00000

```
[5]: df.isnull().sum()
```

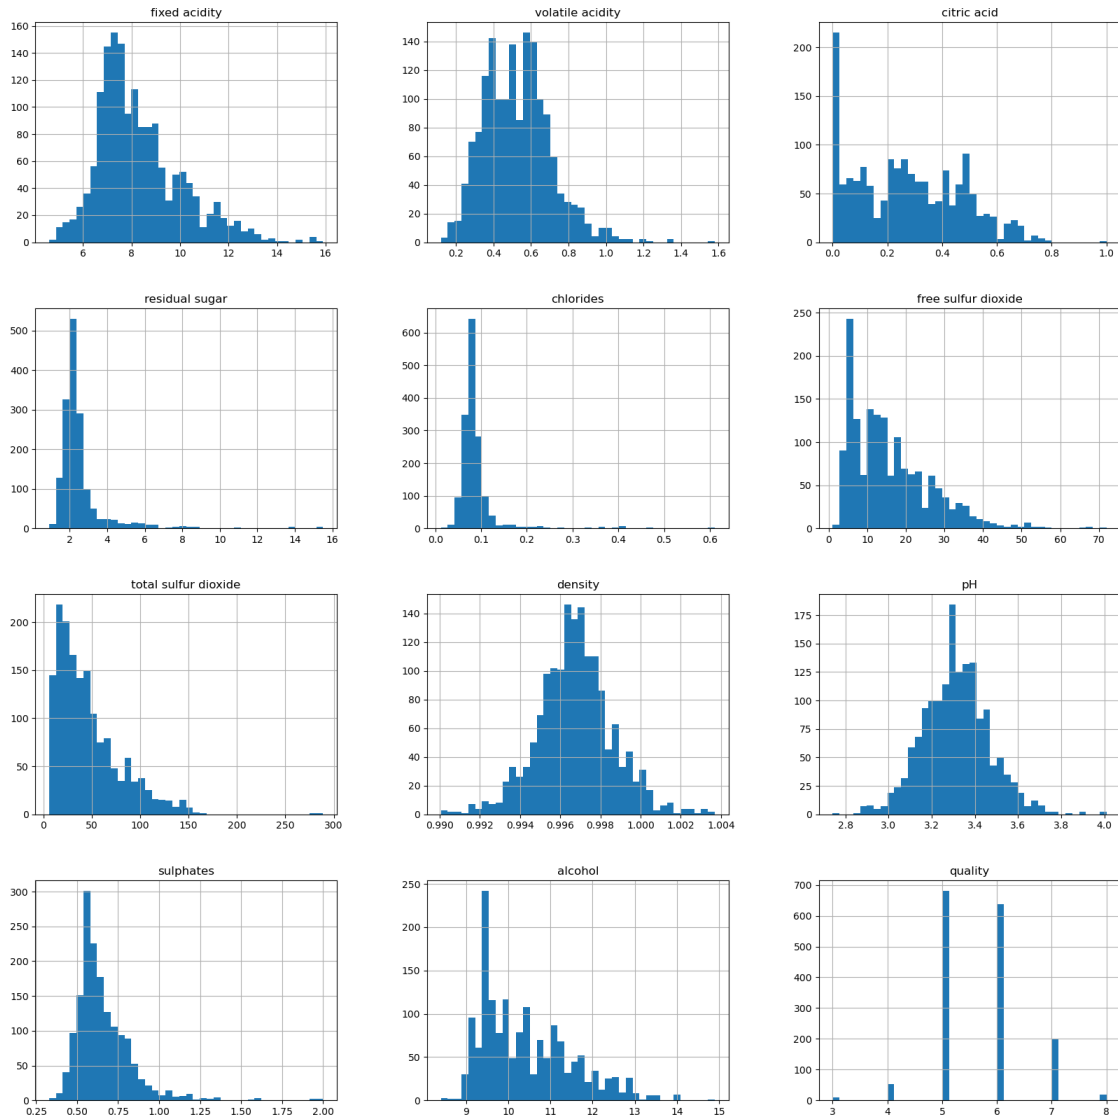
```
[5]: fixed acidity      0
      volatile acidity  0
      citric acid       0
      residual sugar    0
      chlorides         0
      free sulfur dioxide 0
      total sulfur dioxide 0
      density           0
      pH               0
      sulphates         0
      alcohol           0
      quality           0
      dtype: int64
```

```
[9]: for col in df.columns:
      if df[col].isnull().sum() > 0:
          df[col] = df[col].fillna(df[col].mean())

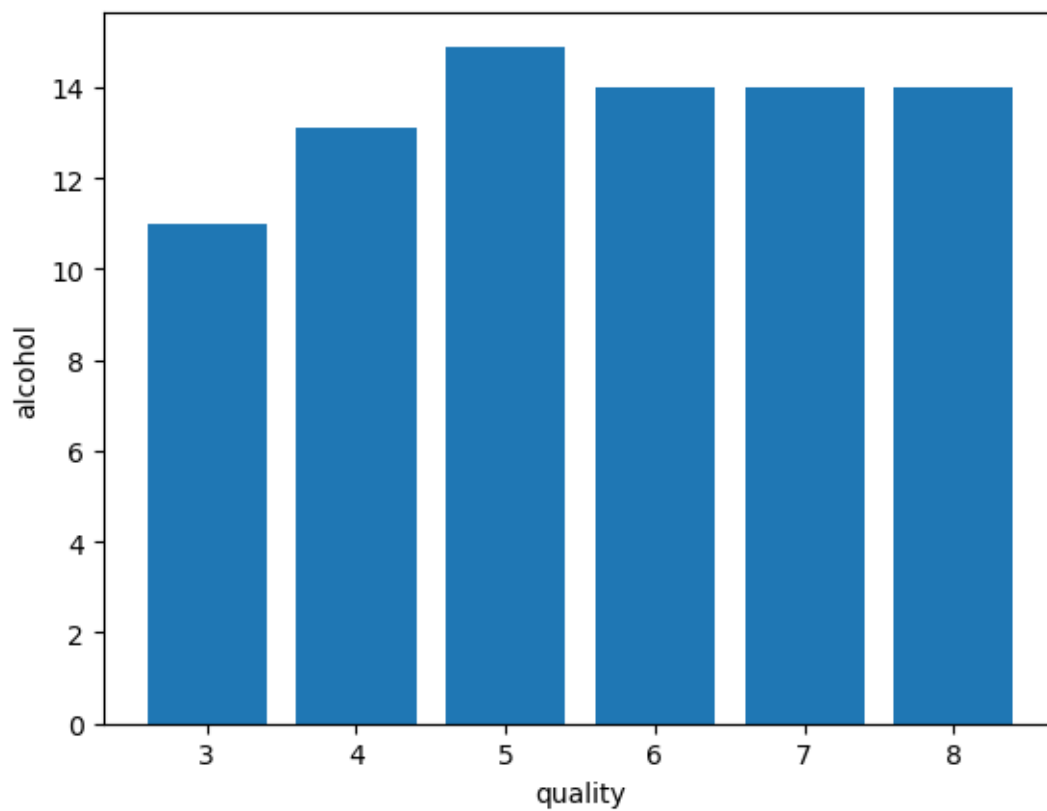
      df.isnull().sum().sum()
```

```
[9]: 0
```

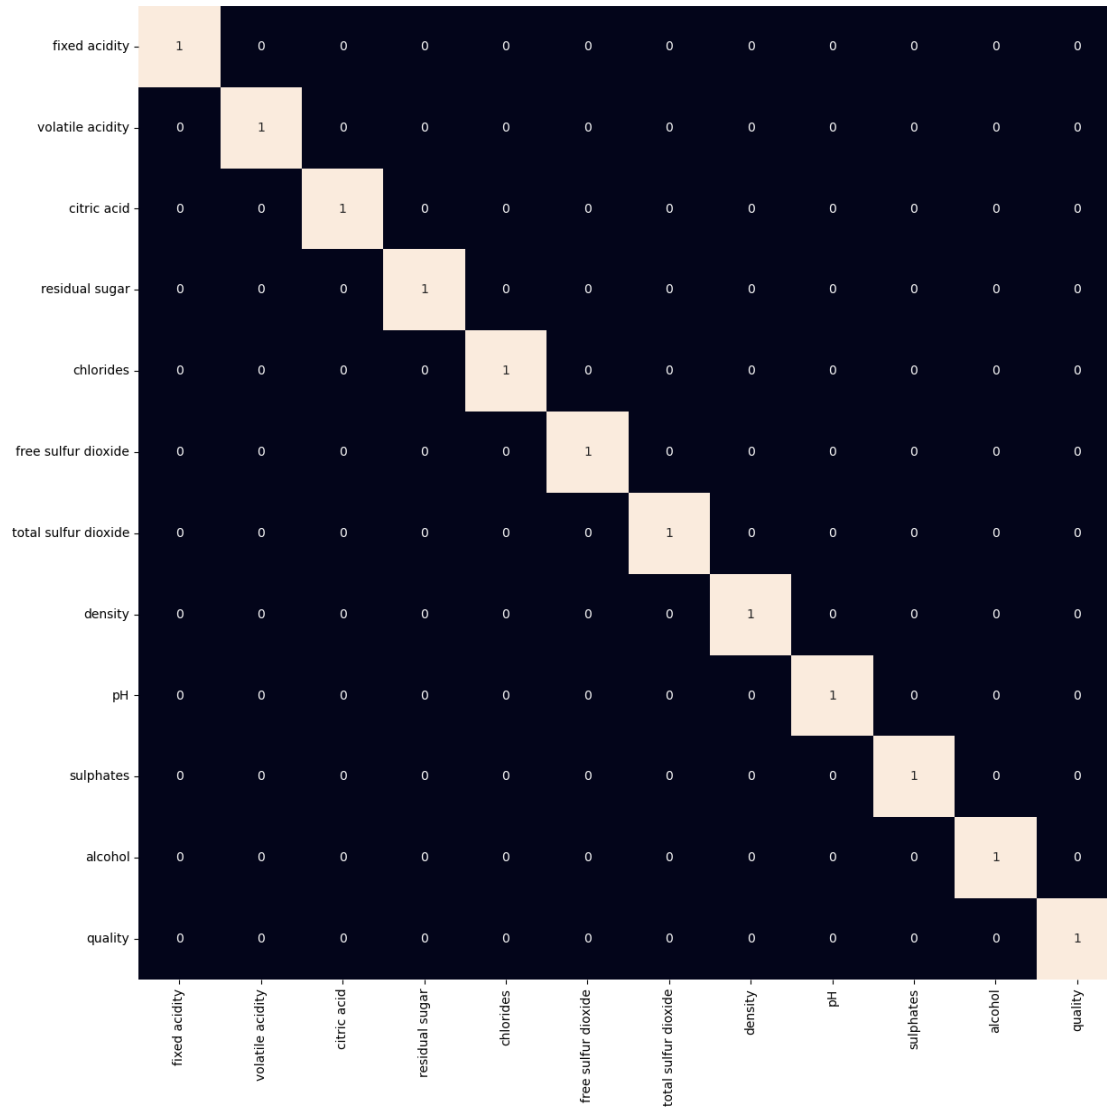
```
[11]: df.hist(bins=40, figsize=(20, 20))
      plt.show()
```



```
[12]: plt.bar(df['quality'], df['alcohol'])
plt.xlabel('quality')
plt.ylabel('alcohol')
plt.show()
```



```
[13]: plt.figure(figsize=(14, 14))  
      sb.heatmap(df.corr() > 0.7, annot=True, cbar=False)  
      plt.show()
```



```
[15]: df = df.drop('total sulfur dioxide', axis=1)
```

```
[16]: df['best quality'] = [1 if x > 5 else 0 for x in df.quality]
```

```
[17]: df.replace({'white': 1, 'red': 0}, inplace=True)
```

```
[18]: features = df.drop(['quality', 'best quality'], axis=1)
target = df['best quality']

xtrain, xtest, ytrain, ytest = train_test_split(features, target, test_size=0.
↪2, random_state=40)

xtrain.shape, xtest.shape
```

```
[18]: ((1279, 10), (320, 10))
```

```
[19]: norm = MinMaxScaler()
      xtrain = norm.fit_transform(xtrain)
      xtest = norm.transform(xtest)
```

```
[36]: models = [LogisticRegression(), XGBClassifier(), SVC(kernel='rbf')]

      for i in range(3):
          models[i].fit(xtrain, ytrain)

          print(f'{models[i]} : ')
          print('Training Accuracy : ', metrics.roc_auc_score(ytrain, models[i].
          ↪predict(xtrain)))
          print('Validation Accuracy : ', metrics.roc_auc_score(ytest, models[i].
          ↪predict(xtest)))
          print()
```

```
LogisticRegression() :
Training Accuracy : 0.7286886534333447
Validation Accuracy : 0.765345444536196
```

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
              predictor=None, random_state=None, ...) :
Training Accuracy : 1.0
Validation Accuracy : 0.8345523180370414
```

```
SVC() :
Training Accuracy : 0.7699408577589806
Validation Accuracy : 0.7930675160237505
```

```
[ ]:
```

```
[38]: print(metrics.classification_report(ytest, models[1].predict(xtest)))
```

	precision	recall	f1-score	support
0	0.81	0.84	0.82	147

1	0.86	0.83	0.84	173
accuracy			0.83	320
macro avg	0.83	0.83	0.83	320
weighted avg	0.84	0.83	0.83	320

[]: